

Human-Participant Ethics and Validation Roadmap

From synthetic authority-allocation experiments to ethically defensible human-AI collaboration studies

Domain-Specific Research Document

Human-AI Codecision Systems Research Program

Stage One revision - September 2026

Evidence status: Synthetic Stage One contains five independent seeds, 180 trials per policy and seed, 3,600 total decisions, and a controlled AI distribution shift in the final quarter. There were zero human participants. Adaptive codecision selectively regulated authority but did not improve synthetic accuracy or declared utility over human-only control. Human trust, satisfaction, and real cognitive state remain unmeasured.

1. Purpose

The current repository contains zero human participants. This roadmap defines what additional evidence and safeguards are required before claims about human trust, workload, uncertainty, satisfaction, reliance, or perceived autonomy are made.

2. Evidence boundary

Construct	Synthetic Stage One	Human validation needed?
Authority allocation	Implemented policy behavior	Yes, for perceived/actual human agency
Accuracy	Synthetic-agent outcome	Yes, for participant performance
Trust	Simulated state only	Yes
Workload	Synthetic feature only	Yes
Satisfaction	Not measured	Yes
Behavioral uncertainty validity	Modeled signals	Yes
Contestability	Implemented mechanism	Yes, for usability/effectiveness

3. Participant-study prerequisites

- Obtain appropriate institutional ethics/IRB determination before recruiting participants when required by the institution and study design.
- Use informed consent describing telemetry, task demands, AI behavior, override/contestability mechanisms, and data retention.
- Collect only data necessary for the stated hypotheses and separate identifying information from analytic telemetry where feasible.
- Do not present simulated trust or workload variables as validated psychological measures.
- Use validated or appropriately justified instruments for trust/workload constructs when those constructs become study outcomes.
- Predefine withdrawal, exclusion, adverse-event, debriefing, and data-handling procedures.
- Audit whether intervention frequency or uncertainty inference differs systematically across participant groups before fairness claims are made.

4. Staged validation

Stage	Goal	Permitted claim
Synthetic	mechanism and threshold validation	policy behavior in simulation
Usability pilot	interface/protocol feasibility	participants can understand/use the system
Approved controlled study	human performance and calibrated reliance	bounded to sampled population/task
External replication	generalizability	only after

		independent/expanded evidence
--	--	-------------------------------

5. High-stakes boundary

The research environment must not be treated as sufficient validation for healthcare, aviation, finance, security operations, or other high-stakes deployment. Each domain introduces separate safety, legal, professional, and accountability requirements. Contestability and explanation are necessary research variables, not proof that override authority is acceptable.